

EFOP-3.6.2-16-2017-00013



European Union

A modern look at GRIN

an optimizing functional language back end

Péter Podlovics, Csaba Hruska

Eötvös Loránd University (ELTE),
Budapest, Hungary

STCS-2019



HUNGARIAN
GOVERNMENT

European Union
European Social
Fund



INVESTING IN YOUR FUTURE

Overview

Introduction

Extensions

Dead Data Elimination

Results

Introduction

Why functional?

- Declarativeness

pro: can program on a higher abstraction level

- Composability

pro: can easily piece together smaller programs

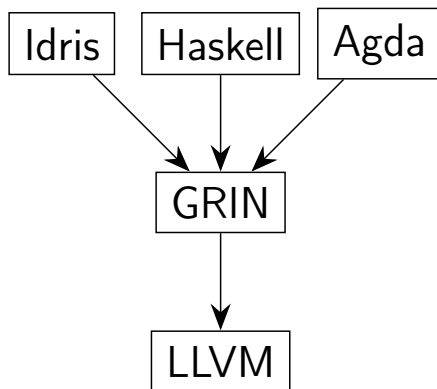
con: results in a lot of function calls

- Functions are first class citizens

pro: higher order functions

con: unknown function calls

Graph Reduction Intermediate Notation



Front end code

```
main = sum (upto 0 10)
```

```
upto n m  
  | n > m = []  
  | otherwise = n : upto (n+1) m
```

```
sum [] = 0
```

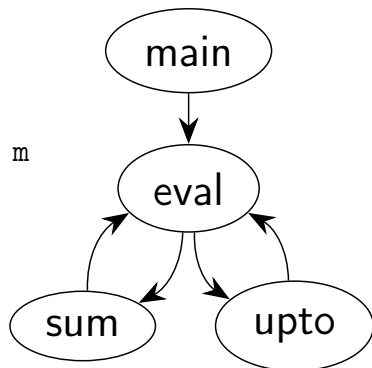
```
sum (x:xs) = x + sum xs
```

Front end code

```
main = sum (upto 0 10)
```

```
upto n m  
  | n > m = []  
  | otherwise = n : upto (n+1) m
```

```
sum [] = 0  
sum (x:xs) = x + sum xs
```



GRIN code

```
grinMain =
```

```
  t1 <- store (CInt 1)
  t2 <- store (CInt 10)
  t3 <- store (Fupto t1 t2)
  t4 <- store (Fsum t3)
  (CInt r) <- eval t4
  _prim_int_print r
```

```
eval p =
  v <- fetch p
  case v of
    (CInt n)      -> pure v
    (CNil)        -> pure v
    (CCons y ys) -> pure v
    (Fupto a b) ->
      zs <- upto a b
      update p zs
      pure zs
    (Fsum c) ->
      s <- sum c
      update p s
      pure s
```


Transformation machinery

- Inline calls to `eval`
- Run dataflow analyses:
 - Heap points-to analysis
 - Sharing analysis
- Run transformations until we reach a fixed-point:
 - Sparse Case Optimization
 - Common Subexpression Elimination
 - Generalized Unboxing
 - etc . . .

Extensions

Extending Heap points-to

1 $\rightarrow \{ \text{CInt}[\{BAS\}] \}$
2 $\rightarrow \{ \text{CInt}[\{BAS\}] \}$
3 $\rightarrow \{ \text{Fupto}[\{1\}, \{2\}], \text{CNil}[], \text{CCons}[\{1, 5\}, \{6\}] \}$
4 $\rightarrow \{ \text{Fsum}[\{3\}], \text{CInt}[\{BAS\}] \}$
5 $\rightarrow \{ \text{CInt}[\{BAS\}] \}$
6 $\rightarrow \{ \text{Fupto}[\{5\}, \{2\}], \text{CNil}[], \text{CCons}[\{1, 5\}, \{6\}] \}$

Extending Heap points-to

1 $\rightarrow \{ \text{CInt}[\{BAS\}] \}$
2 $\rightarrow \{ \text{CInt}[\{BAS\}] \}$
3 $\rightarrow \{ \text{Fupto}[\{1\}, \{2\}], \text{CNil}[], \text{CCons}[\{1, 5\}, \{6\}] \}$
4 $\rightarrow \{ \text{Fsum}[\{3\}], \text{CInt}[\{BAS\}] \}$
5 $\rightarrow \{ \text{CInt}[\{BAS\}] \}$
6 $\rightarrow \{ \text{Fupto}[\{5\}, \{2\}], \text{CNil}[], \text{CCons}[\{1, 5\}, \{6\}] \}$

$BAS \in \{\text{Int64}, \text{Float}, \text{Bool}, \text{String}, \text{Char}\}$

Extending Heap points-to

$1 \rightarrow \{ \text{CInt}[\{BAS\}] \}$
 $2 \rightarrow \{ \text{CInt}[\{BAS\}] \}$
 $3 \rightarrow \{ \text{Fupto}[\{1\}, \{2\}], \text{CNil}[], \text{CCons}[\{1, 5\}, \{6\}] \}$
 $4 \rightarrow \{ \text{Fsum}[\{3\}], \text{CInt}[\{BAS\}] \}$
 $5 \rightarrow \{ \text{CInt}[\{BAS\}] \}$
 $6 \rightarrow \{ \text{Fupto}[\{5\}, \{2\}], \text{CNil}[], \text{CCons}[\{1, 5\}, \{6\}] \}$

$BAS \in \{\text{Int64}, \text{Float}, \text{Bool}, \text{String}, \text{Char}\}$

`indexArray# :: Array# a -> Int# -> (# a #)`
`newMutVar# :: a -> s -> (# s, MutVar# s a #)`

LLVM back end

```
grinMain =  
  t1 <- store (CInt 1)  
  t2 <- store (CInt 10)  
  t3 <- store (Fupto t1 t2)  
  t4 <- store (Fsum t3)  
  (CInt r') <- eval t4  
  _prim_int_print r'
```

```
upto m n =  
  (CInt m') <- eval m  
  (CInt n') <- eval n  
  b' <- _prim_int_gt m' n'  
  case b' of  
    #True -> pure (CNil)
```

```
sum l = ...
```

```
eval p = ...
```

LLVM back end

```
grinMain =  
  t1 <- store (CInt 1)  
  t2 <- store (CInt 10)  
  t3 <- store (Fupto t1 t2)  
  t4 <- store (Fsum t3)  
  (CInt r') <- eval t4  
  _prim_int_print r'
```

```
upto m n =  
  (CInt m') <- eval m  
  (CInt n') <- eval n  
  b' <- _prim_int_gt m' n'  
  case b' of  
    #True -> pure (CNil)
```

```
sum l = ...
```

```
eval p = ...
```

```
grinMain =  
  n1 <- sum 0 1 10  
  _prim_int_print n1  
  
sum s lo hi =  
  b <- _prim_int_gt lo hi  
  if b then  
    pure s  
  else  
    lo' <- _prim_int_add lo 1  
    s' <- _prim_int_add s lo  
    sum s' lo' hi
```

LLVM back end

```
grinMain =  
  t1 <- store (CInt 1)  
  t2 <- store (CInt 10)  
  t3 <- store (Fupto t1 t2)  
  t4 <- store (Fsum t3)  
  (CInt r') <- eval t4  
  _prim_int_print r'
```

```
upto m n =  
  (CInt m') <- eval m  
  (CInt n') <- eval n  
  b' <- _prim_int_gt m' n'  
  case b' of  
    #True -> pure (CNil)
```

```
sum l = ...
```

```
eval p = ...
```

```
grinMain =  
  n1 <- sum 0 1 10  
  _prim_int_print n1  
  
sum s lo hi =  
  b <- _prim_int_gt lo hi  
  if b then  
    pure s  
  else  
    lo' <- _prim_int_add lo 1  
    s' <- _prim_int_add s lo  
    sum s' lo' hi
```

```
grinMain:  
# BB#0:  
movabsq    $55, %rdi  
jmp        _prim_int_print
```


Dead Data Elimination

Dead data elimination

```
length : List a -> Nat
length Nil = Z
length (Cons x xs)
  = S (length xs)
```

$\xRightarrow{\text{DDE}}$

```
length p =
  xs <- fetch p
  case xs of
    (Cons ys) ->
      l1 <- length ys
      l2 <- _prim_int_add l1 1
      pure l2
    (Nil) ->
      pure 0
```

Applications

- $\text{Map} \rightarrow \text{Set}$
- Type class dictionaries
- Type erasure for dependently typed languages

What do we need?

- Producers & consumers
- Detect dead fields
- Connect consumers to producer
- Remove or transform dead fields

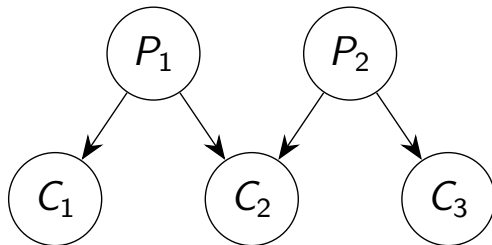
```
grinMain =
```

```
  a0 <- pure 5
  n0 <- pure (CNil)
  p0 <- store n0
  n1 <- pure (CCons a0 p0)
  r  <- case n1 of
    (CNil)  ->
      pure (CNil)
    (CCons x xs) ->
      xs' <- fetch xs
      pure xs'
  pure r
```

Producers

```
a0    -> {}
n0    -> {CNil{n0}}
n1    -> {CCons{n1}}
p0    -> {}
r     -> {CNil{n0}}
x     -> {}
xs    -> {}
xs'   -> {CNil{n0}}
```

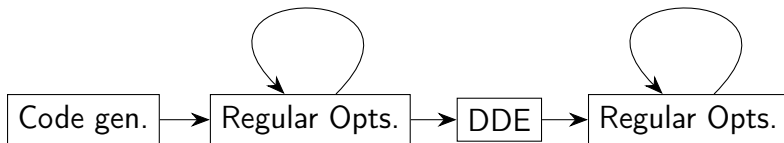
Producers and consumers



Results

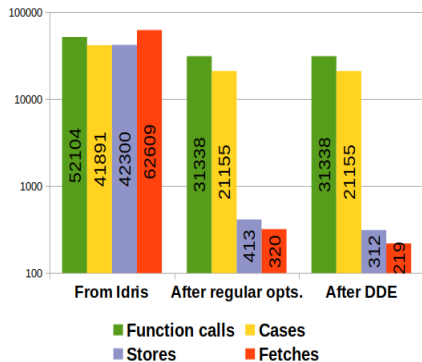
Setup

- Small Idris code snippets from:
Type-driven Development with Idris by Edwin Brady
- Only interpreted code
- Compile- & runtime measurements
- Pipeline setup:

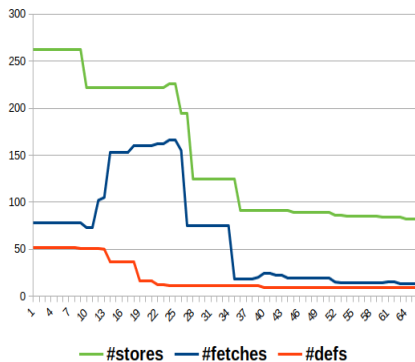


Length

Runtime Statistics

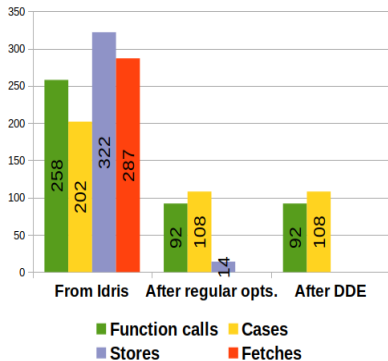


Compile Time Statistics

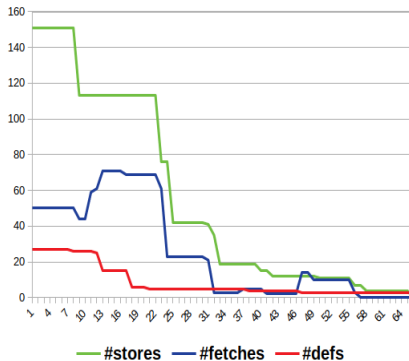


Exact length

Runtime Statistics

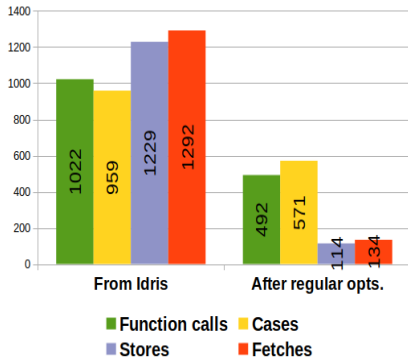


Compile Time Statistics




Reverse

Runtime Statistics

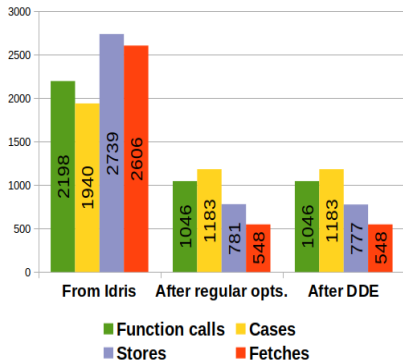


Compile Time Statistics

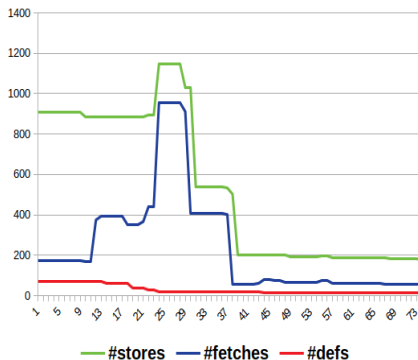


Type level functions

Runtime Statistics



Compile Time Statistics



Conclusions

- The optimizer works well:
 - the number of stores, fetches, function calls and pattern matches significantly decreased
 - the structure of the code resembles that of an imperative language
- Dead Data Elimination:
 - is a bit costly
 - is a specific optimization
 - can completely transform data structures
 - can trigger further transformations

EFOP-3.6.2-16-2017-00013



European Union

THANK YOU FOR YOUR ATTENTION!

SZÉCHENYI 2020



HUNGARIAN
GOVERNMENT

European Union
European Social
Fund



INVESTING IN YOUR FUTURE

Sparse case optimization

<m0>

v <- eval l

case v of

CNil -> <m1>

CCons x xs -> <m2>

$v \in \{\text{CCons}\}$
 \Longrightarrow

<m0>

v <- eval l

case v of

CCons x xs -> <m2>

Compiled data flow analysis

- Analyzing the syntax tree has an interpretation overhead
- We can work around this by "compiling" our analysis into an executable program
- The compiled abstract program is independent of the AST
- It can be executed in a different context (ie.: by another program or on GPU)
- After run (iteratively), it produces the result of the given analysis

A small functional program

```
main = sum (upto 0 10)
```

```
upto from to  
  | from > to = []  
  | otherwise = from : upto (from+1) to
```

```
sum [] = 0
```

```
sum (x:xs) = x + sum xs
```

Strict control flow

```
main = sum (upto 0 10)
```

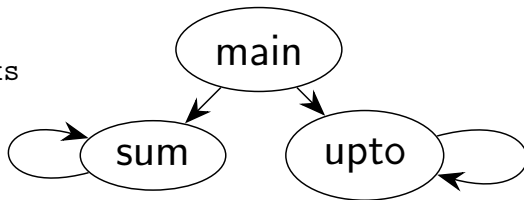
```
upto m n
```

```
  | m > n = []
```

```
  | otherwise = m : upto (m+1) n
```

```
sum [] = 0
```

```
sum (x:xs) = x + sum xs
```



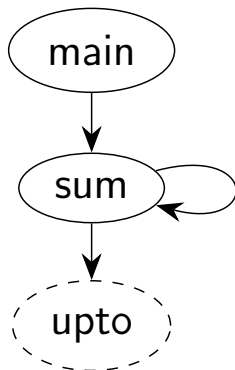
Optimized lazy control flow

```
main = sum (upto 0 10)
```

```
upto m n  
  | m > n = []  
  | otherwise = m : upto (m+1) n
```

```
sum [] = 0
```

```
sum (x:xs) = x + sum xs
```



Goals

- We need to handle laziness
- We need to optimize across functions
- Accomplish both of these for all functional languages

Properties

- Designed for the computer
- Simple syntax, and semantics
- Untyped, but we use a typed version (for LLVM)
- First order language
- Monadic structure
- Single Static Assignment property
- Explicit laziness
- Global `eval` (generated)
- No unknown function calls

Semantics

- C, F, P nodes
- Only basic values and pointers can be in nodes
- Functions cannot return pointers
 - More register usage is exposed
 - The caller can decide whether the return value should be put onto the heap
- `store`, `fetch`, `update`
- Control flow can only diverge and merge at case expressions

Laziness in GRIN

```
upto m n =  
  (CInt m') <- eval m  
  (CInt n') <- eval n  
  b' <- _prim_int_gt m' n'  
  if b' then  
    pure (CNil)  
  else  
    m1' <- _prim_int_add m' 1  
    m1 <- store (CInt m1')  
    p <- store (Fupto m1 n)  
    pure (CCons m p)
```


Dead data elimination

<m0>

```
n <- pure (CPair a b)
(CPair x y) <- pure n
```

<m1>

$\xRightarrow{x \text{ is dead}}$

<m0>

```
n <- pure (CPair b)
(CPair y) <- pure n
```

<m1>

Analysis types

- Whole program analysis

The entire program is subject to the analysis

- Interprocedural program analysis

The analysis is performed across functions

- Context insensitive program analysis

- Information is not propagated back to the call site

Heap-points-to

```
grinMain =
```

```
  a0 <- pure 5
```

```
  n0 <- pure (CNil)
```

```
  p0 <- store n0
```

```
  n1 <- pure (CCons a0 p0)
```

```
  r <- case n1 of
```

```
    (CNil) ->
```

```
      pure (CNil)
```

```
    (CCons x xs) ->
```

```
      xs' <- fetch xs
```

```
      pure xs'
```

```
  pure r
```

Heap

```
0 -> {CNil[]}
```

Env

```
a0 -> {T_Int64}
```

```
n0 -> {CNil[]}
```

```
n1 -> {CCons[{T_Int64},{0}]}
```

```
p0 -> {0}
```

```
r -> {CNil[]}
```

```
x -> {T_Int64}
```

```
xs -> {0}
```

```
xs' -> {CNil[]}
```

Function

```
grinMain :: {CNil[]}
```

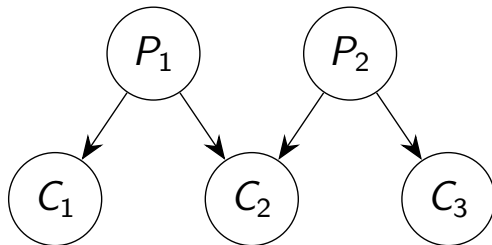
```
grinMain =
```

```
  a0 <- pure 5
  n0 <- pure (CNil)
  p0 <- store n0
  n1 <- pure (CCons a0 p0)
  r  <- case n1 of
    (CNil)  ->
      pure (CNil)
    (CCons x xs) ->
      xs' <- fetch xs
      pure xs'
  pure r
```

Producers

```
a0    -> {}
n0    -> {CNil{n0}}
n1    -> {CCons{n1}}
p0    -> {}
r     -> {CNil{n0}}
x     -> {}
xs    -> {}
xs'   -> {CNil{n0}}
```

Producers and consumers



Liveness

```
grinMain =  
  a0 <- pure 5  
  n0 <- pure (CNil)  
  p0 <- store n0  
  n1 <- pure (CCons a0 p0)  
  r <- case n1 of  
    (CNil) ->  
      pure (CNil)  
    (CCons x xs) ->  
      xs' <- fetch xs  
      pure xs'  
  pure r
```

Heap

0 -> {CNil[]}

Env

a0 -> DEAD

n0 -> {CNil[]}

n1 -> {CCons[DEAD, LIVE]}

p0 -> LIVE

r -> {CNil[]}

x -> DEAD

xs -> LIVE

xs' -> {CNil[]}

Function

grinMain :: {CNil[]}

Results

- eval inlining impact on code size
- dead code elimination impact on code size
- dead code elimination impact on performance
- comparing intra- and interprocedural dead code elimination
- how costly they are?
- how the resulting codes differ?
- how should the transformations be ordered to minimize compilation time, and maximize performance?
- how costly are the analyses?
- how does the GRIN optimized code compare to GHC's?

Summary

- Compiling functional programs has its own challenges
- We can make it easier by introducing a new IR
- We can perform elaborate dataflow analyses on the IR, then ...
- By transforming the code to a more manageable format, we can utilize the already existing infrastructure of LLVM